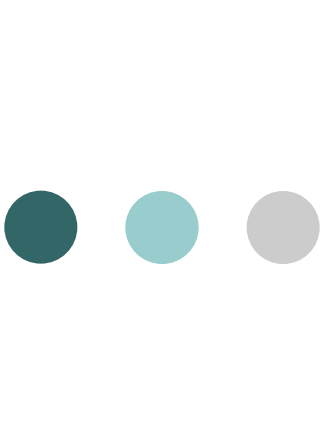

Implementing a DB Server (V1.13)

Operating Systems

Computer Degree

Depto. de Arquitectura de Computadores
Universidad de Málaga

Sections

- DBMS Internal Structure
- Internal APIs

Contents

DBMS INTERNAL STRUCTURE

DBMS Server

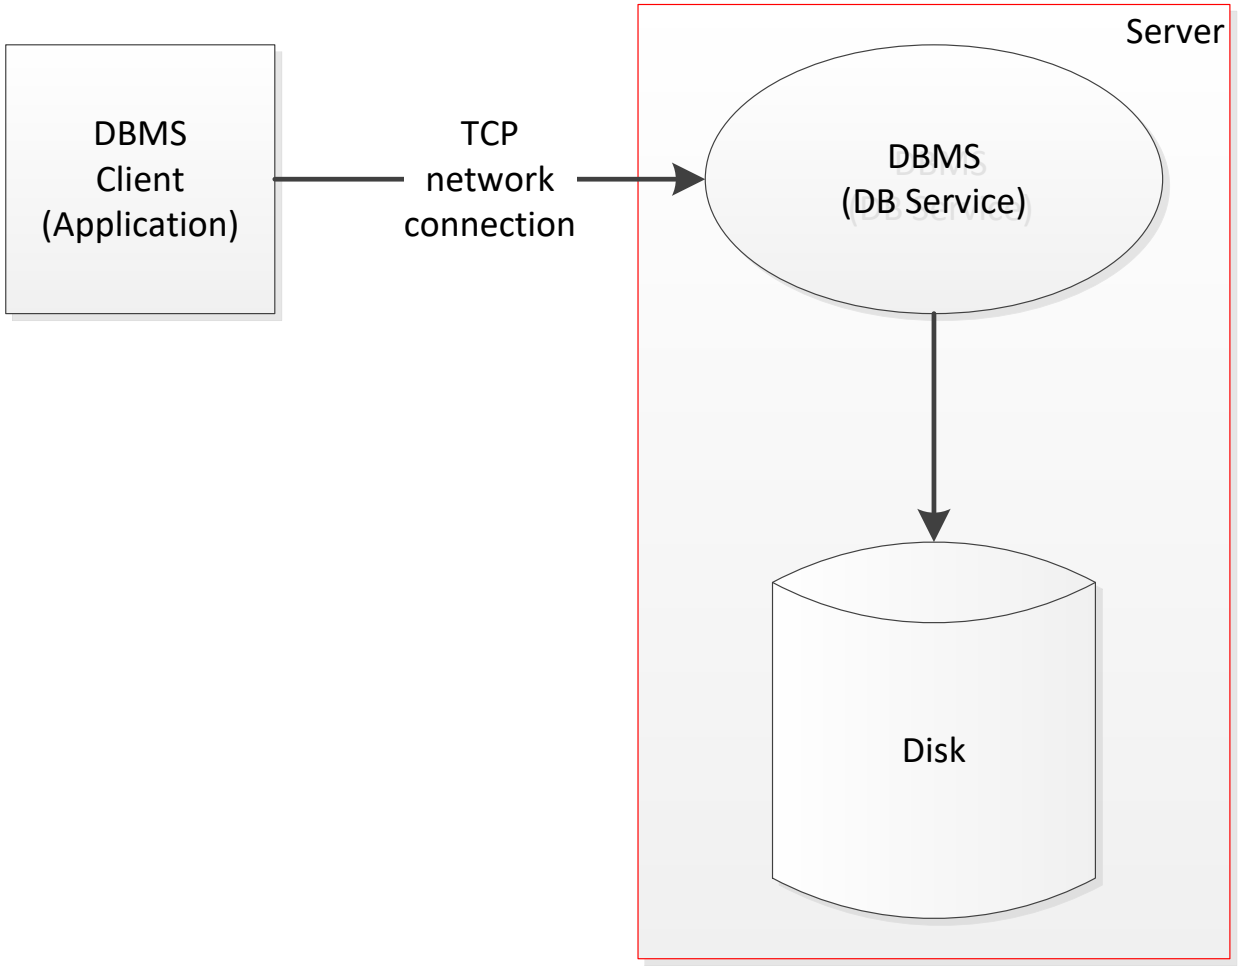
DBMS
Client
(Application)
TCP
network
connection
Server
DBMS
(DB Service)
Disk

DBMS Internal Tasks

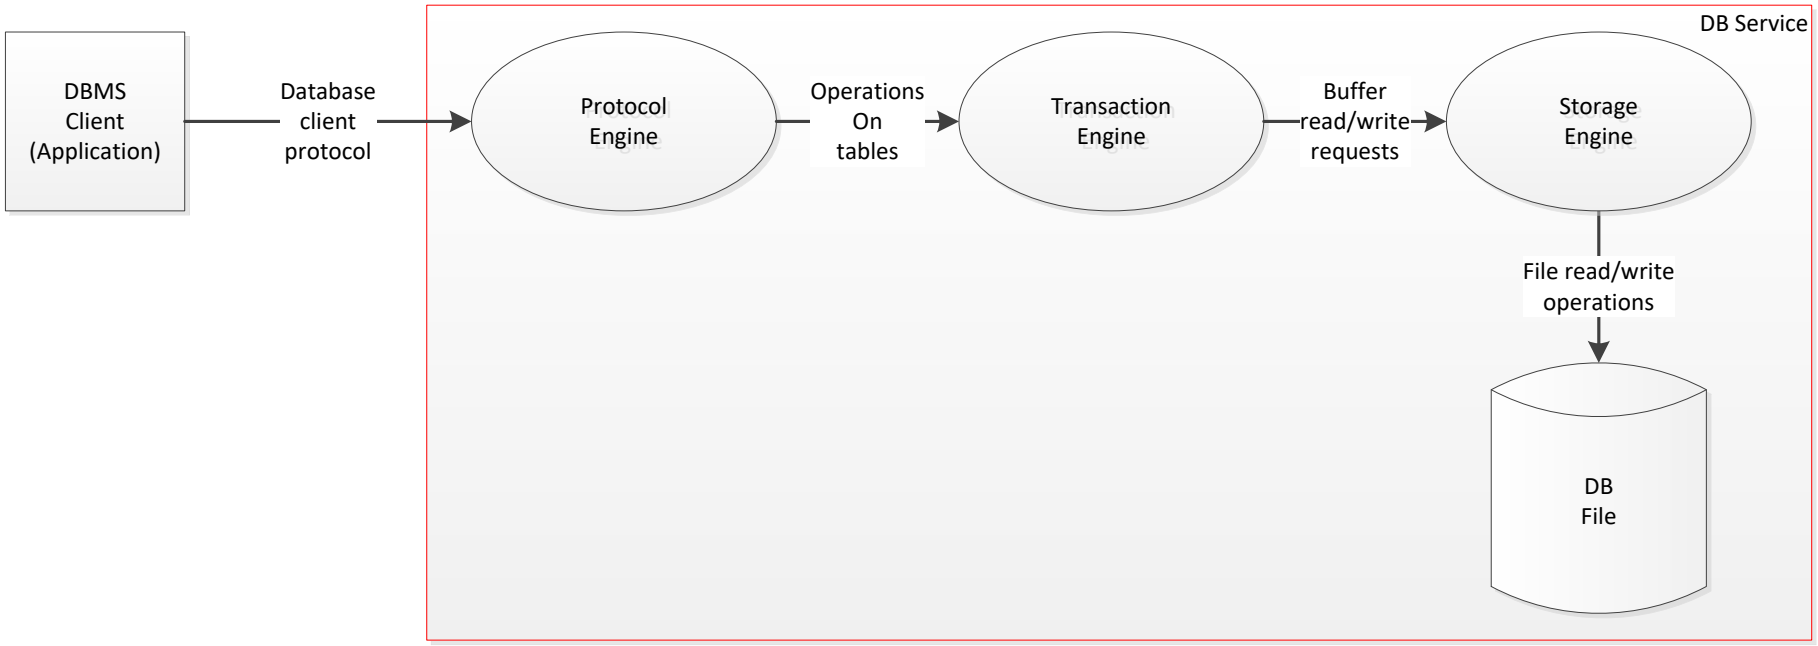

Protocol Engine: Connection Listener

- Connection listener only listens on a TCP port for new connections and creates one session manager for each connection.
- Each session manager decodes SQL syntax and converts commands into operations on tables for only one client.

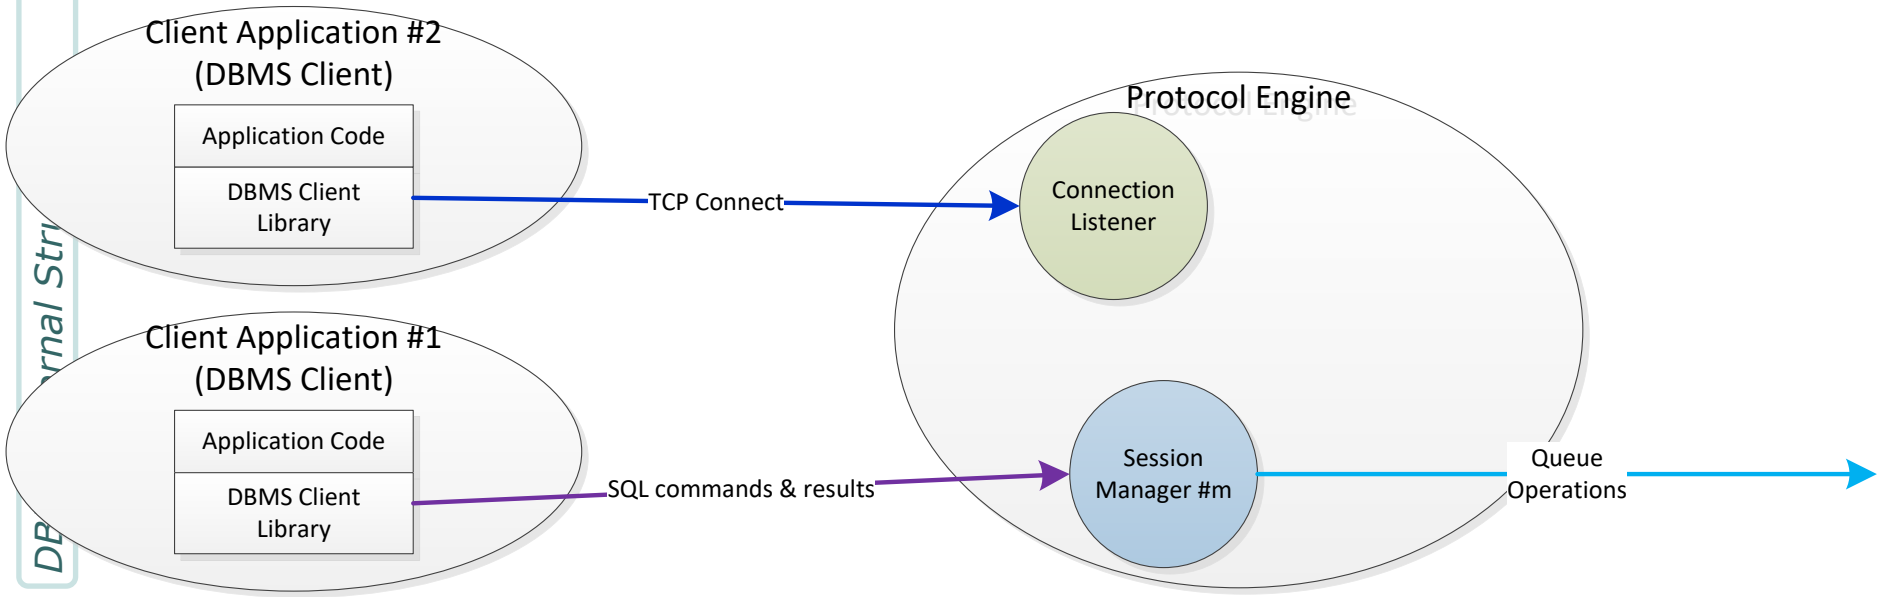

Protocol Engine: Session Manager

- Each client application dialogs with its own session manager process.

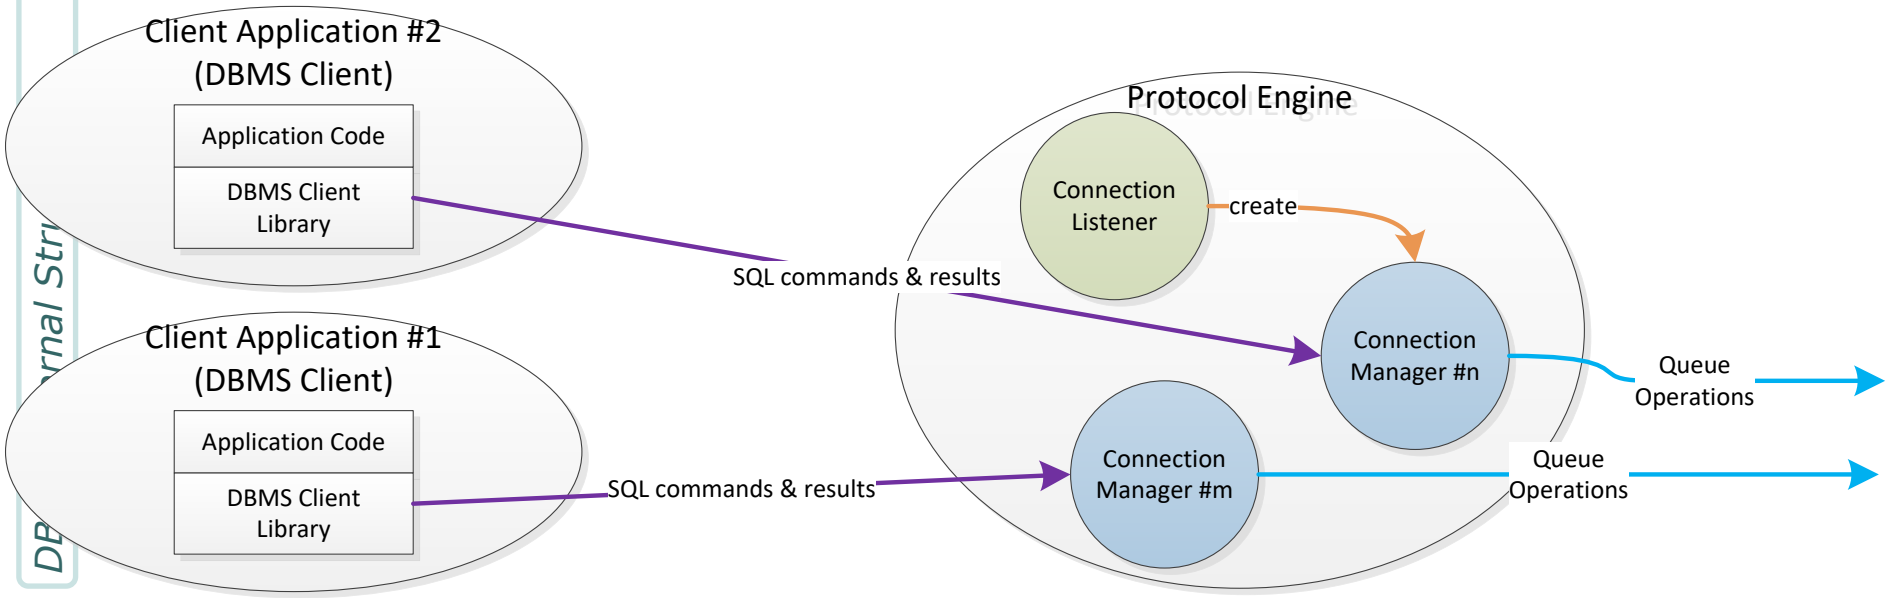

Transaction Engine

- Executes each operation in a new thread to achieve as much parallelism as possible.
- This engine also controls locks on registers or tables.

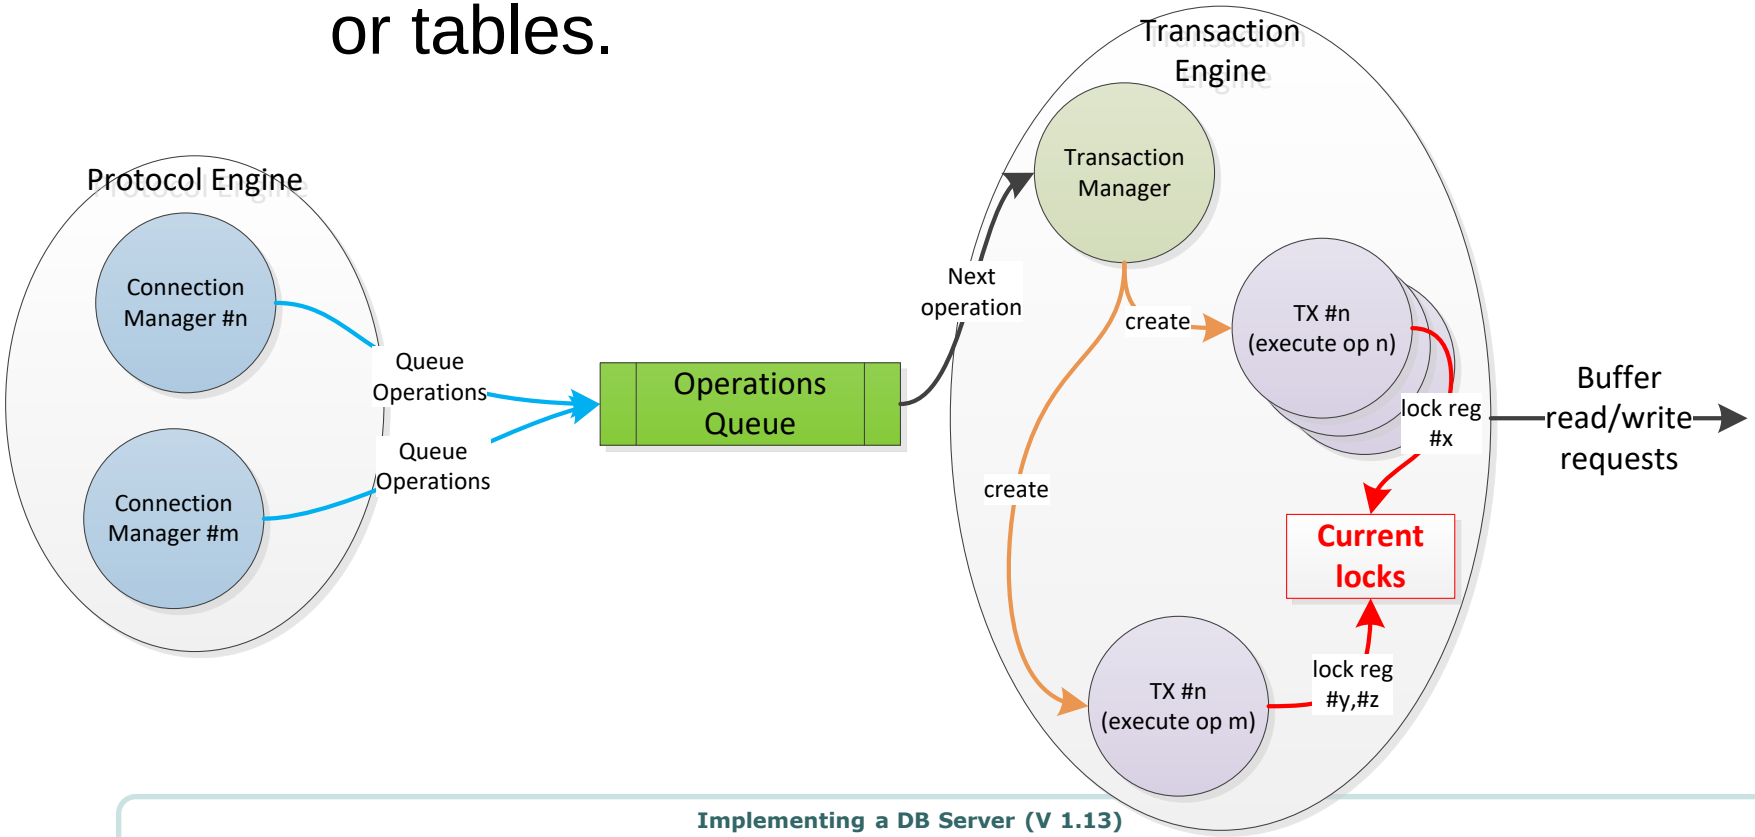

Storage Engine

- Maintains a cache of the file in memory.
- Reads/writes blocks of cache from/to disk in parallel using dedicated threads.

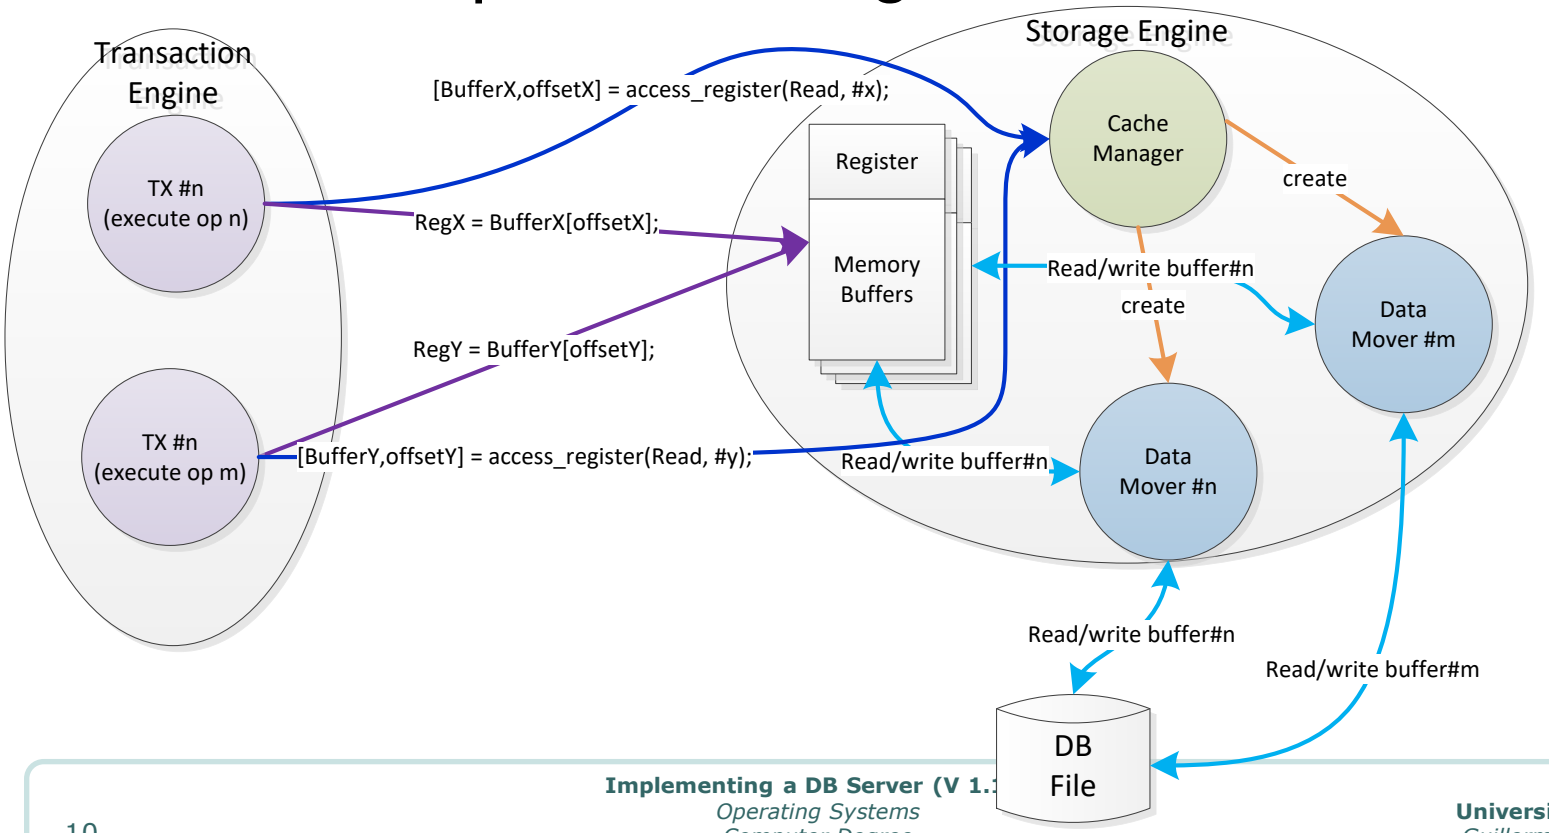

Protocol Engine: Session Manager

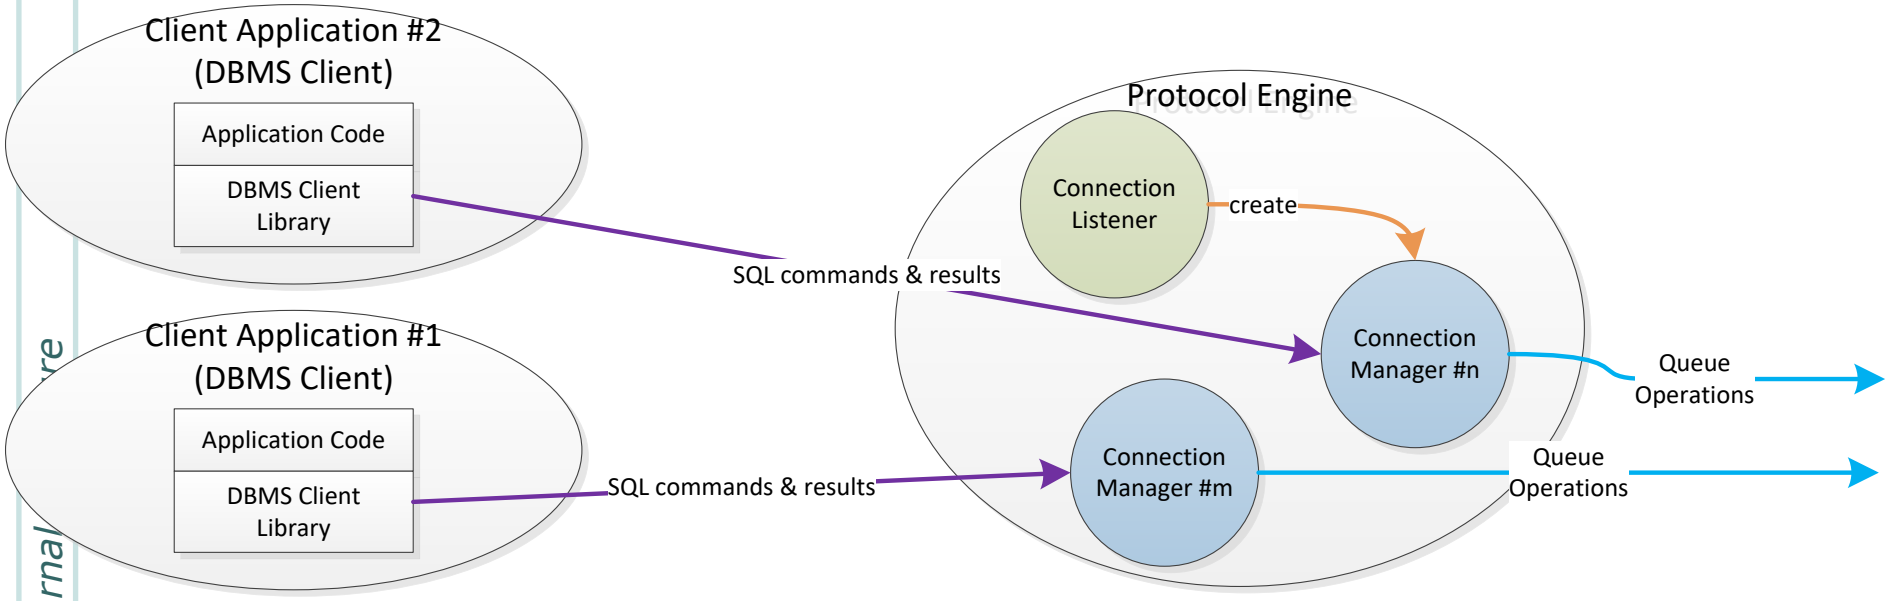

Transaction Engine

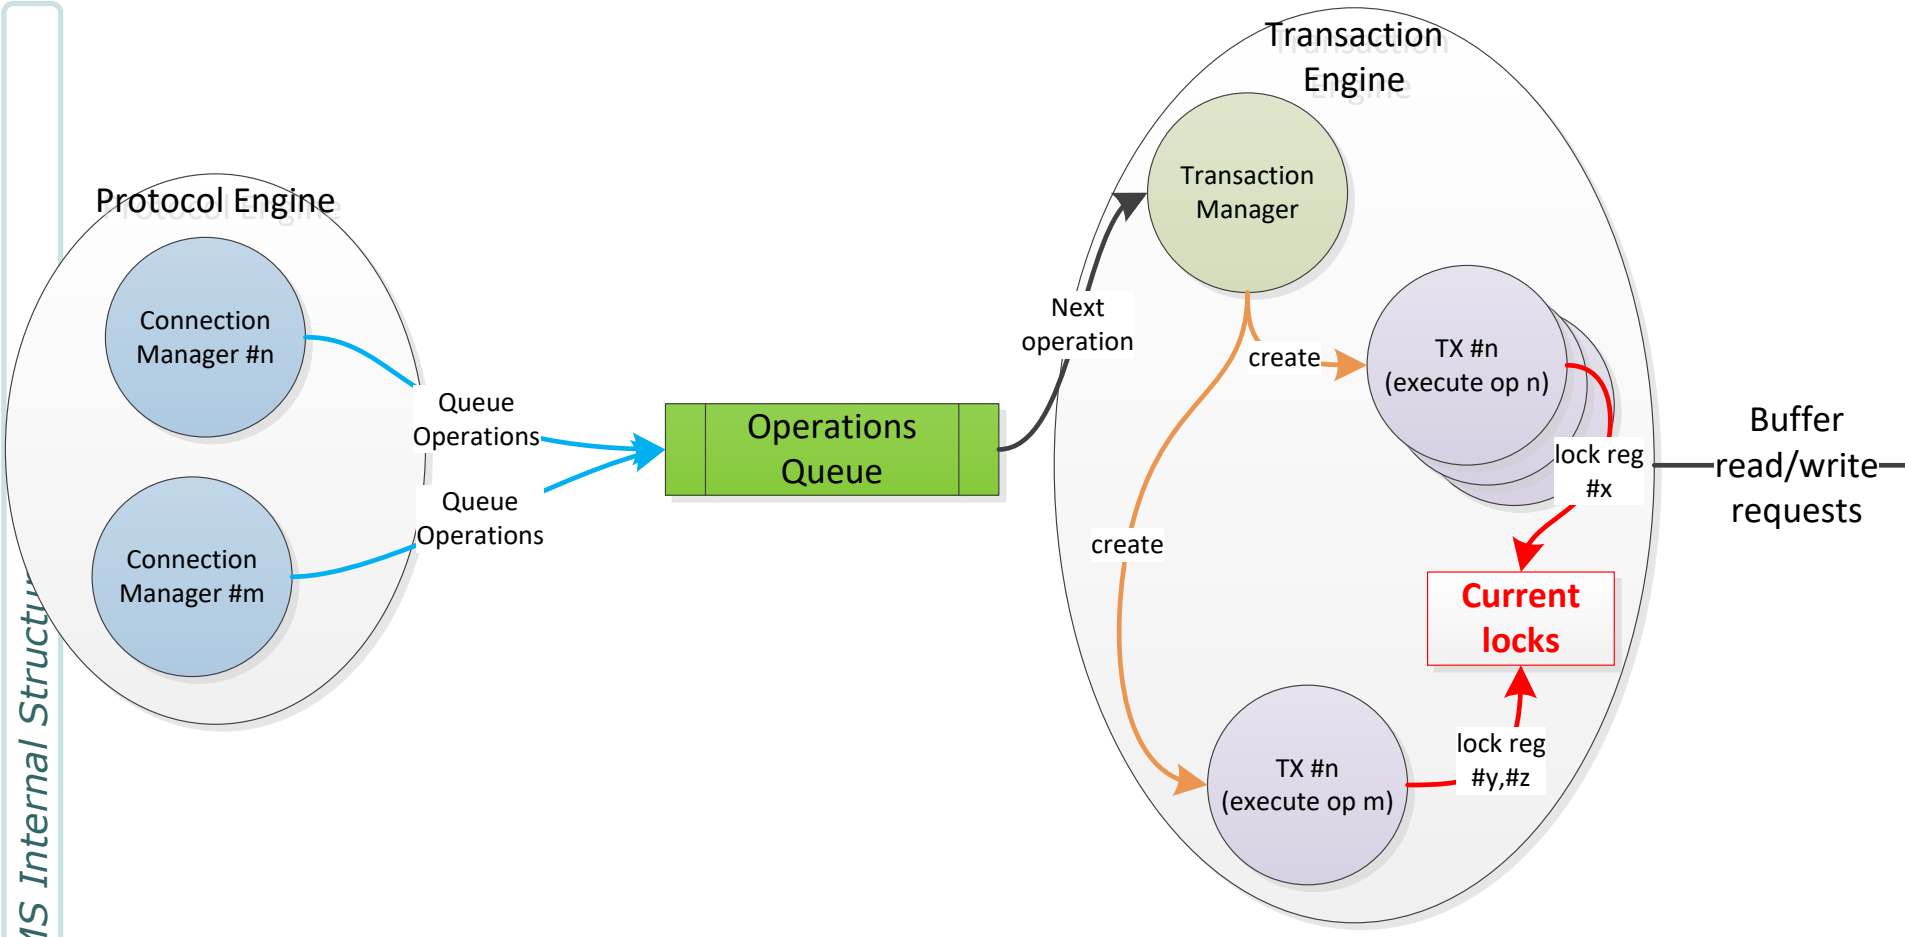

Storage Engine

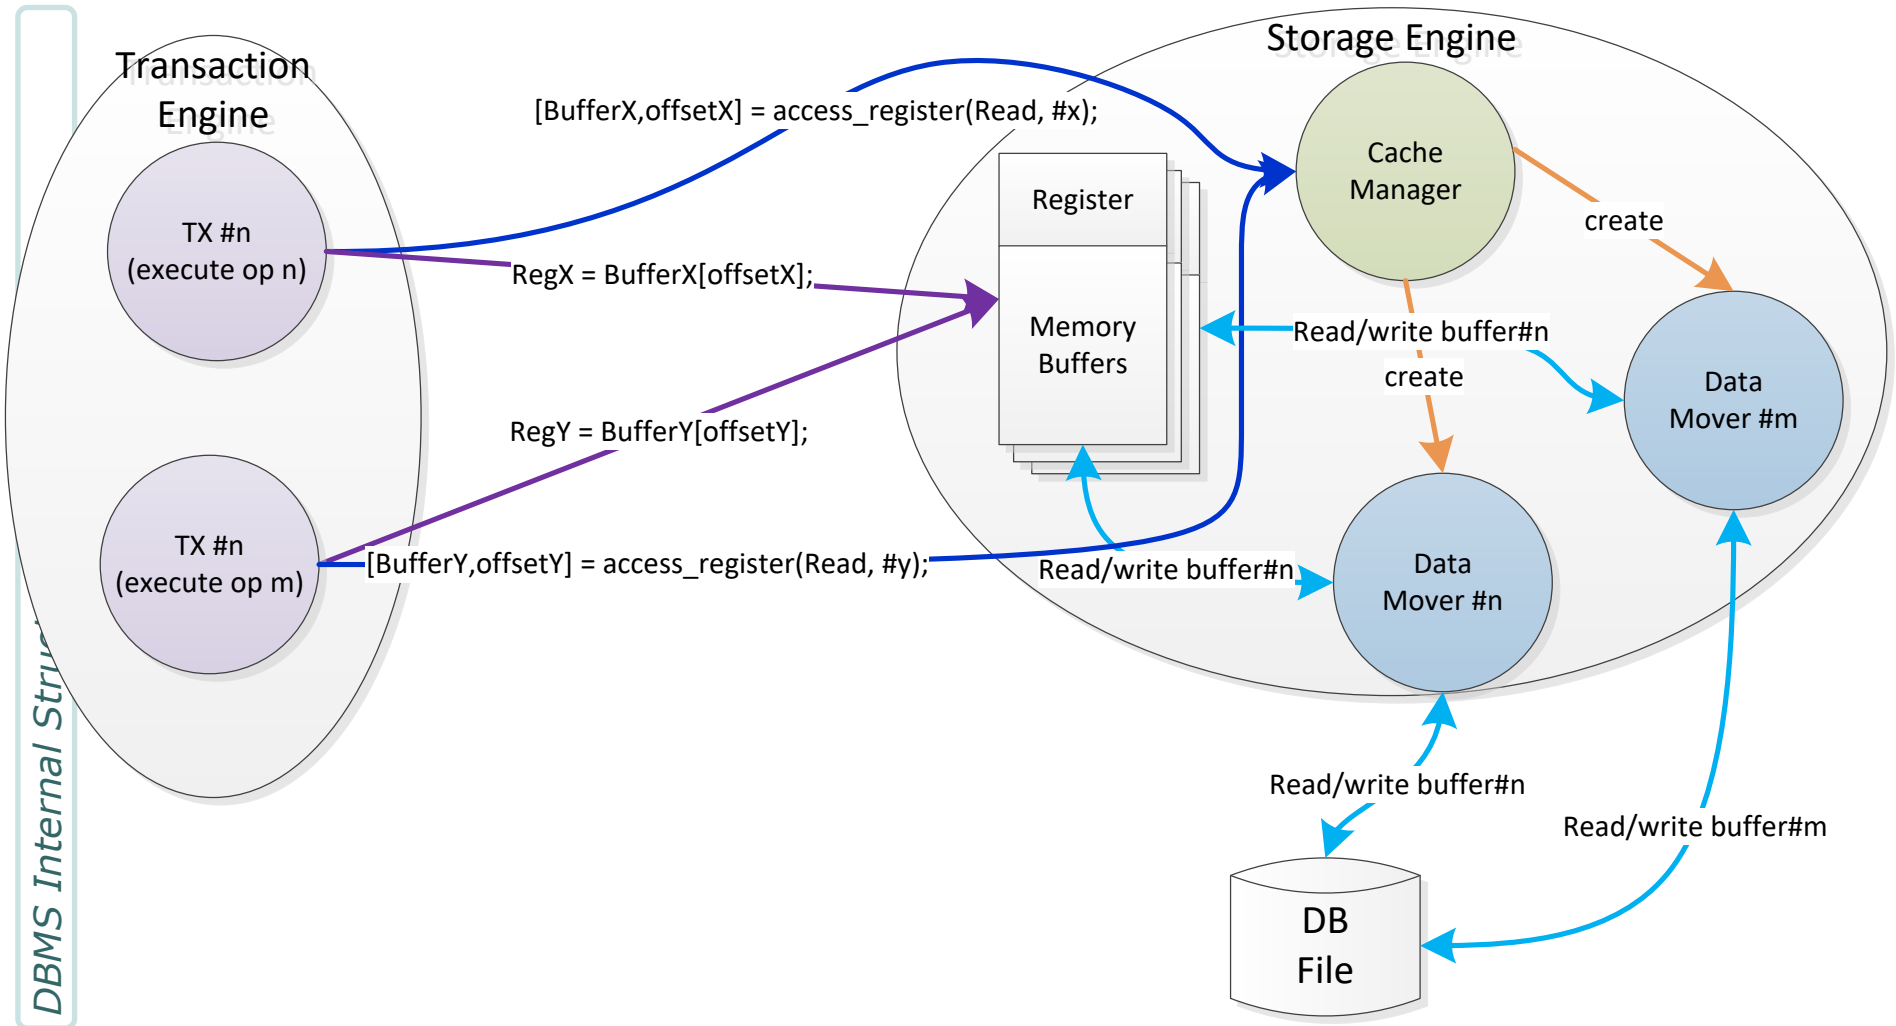

INTERNAL APIS

Need for internal APIs

- Each layer (engine) needs to provide an API for its client layer.
  - The API is a **client library** which can be used to program the upper layer.
    - It only contains the minimal set of functions needed to dialog with the engine.
  - The implementation of the API is done by the lower layer.
  - The interaction between the client library and the lower layer is implemented using **communication services** of the OS like:
    - Semaphores
    - Message queues
    - Shared memory

How to implement an API on OS

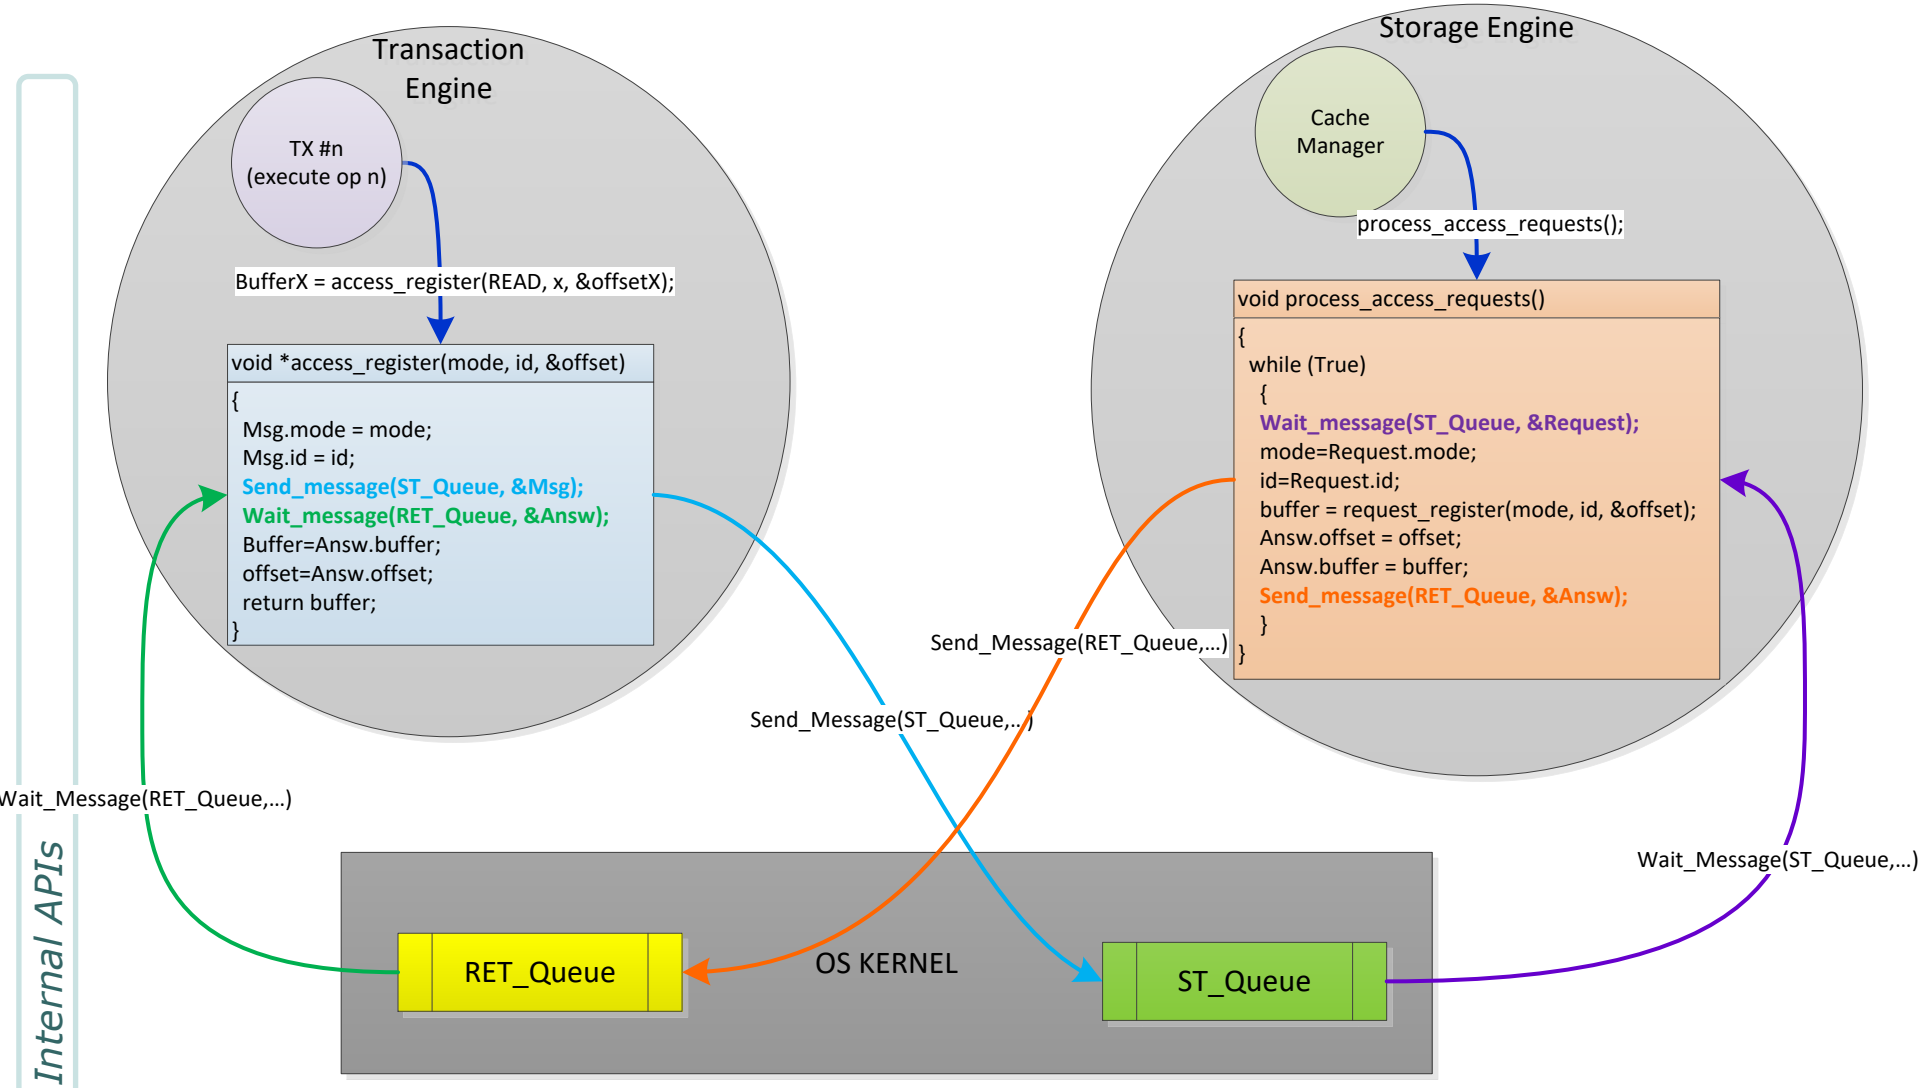

Internal APIs

Proposed functions for APIs

- Protocol engine API (DB protocol)
  - Used to write client applications

```
struct Connection *con;
int status;
char *sql_cmd="select * from T;"
struct Results *results_ptr;

con = connect(serverip);
status = execute_SQL(sql_cmd, &results_ptr);
disconnect(con);
```

Proposed functions for APIs

- Transaction engine API
  - Used to write the session manager of Protocol Engine.

```
struct tQueue *tx_queue, *tx_queue;
Struct tOperation *operation;
struct tResults *result;

tx_queue = get_opqueue();
rx_queue = new_rxqueue(sessionid);
queue_operation(tx_queue, operation, rx_queue);
wait_result(rx_queue, &result);
delete_rxqueue(rx_queue);
```

Proposed functions for APIs

- Storage engine API
  - Used to write the execution thread in Transaction Engine.

```
tMode mode;
int regid;
struct Register *buffer;
long offset;

mode = REG_WRITE;
mode = REG_WRITESYNC;
mode = REG_READ;
status = access_register(mode, regid, &buffer, &offset);
buffer[offset] = register;
status = release_register(regid);
```